
CSCI 3388 Practice Midterm **SOLUTIONS****Instructions:**

- **DO NOT DETACH ANY PAGES—EVEN THE SCRATCH PAGES—FROM THIS EXAM!** Removing page(s) may result in a deduction in your grade.
- This exam is closed book and closed notebook, with no electronic devices allowed. The only resource you may use is one 8.5×11 -inch double-sided study sheet that you prepared yourself.
- You are not allowed to discuss the exam with anyone until the solutions are published.

Any deviation from these rules will constitute a violation of the University policies on academic integrity.

- The exam is designed for 75 minutes and consists of multiple-choice and written-answer problems.
- For the multiple-choice problems, mark your answers by filling in the bubble to the left of your choice. Select only one choice.
- For the written-answer problems, write your answers clearly in the provided space.
- The exam has 6 pages printed on both sides, including this page and three scratch pages at the end.

Name: _____

Email: _____

Eagle ID: _____

Section	Points
1-4 (Multiple Choice)	/ 20
5-8 (Written Answer)	/ 80
Total	/ 100

Multiple Choice (5 points each)

1. N-gram Language Models

Given the following bigram counts from a corpus:

- $\text{count}(\text{"the"}, \text{"cat"}) = 15$
- $\text{count}(\text{"the"}) = 100$
- $\text{count}(\text{"cat"}) = 20$

What is $P(\text{"cat"} \mid \text{"the"})$ using maximum likelihood estimation?

- ☐ 0.15
- ☒ **0.20**
- ☐ 0.75
- ☐ 0.80

Solution: (A). $P(\text{"cat"} \mid \text{"the"}) = \text{count}(\text{"the"}, \text{"cat"}) / \text{count}(\text{"the"}) = 15/100 = 0.15$

2. Neural Network Activation Functions

Which activation function is most commonly used in modern deep learning to address the vanishing gradient problem?

- ☐ Sigmoid
- ☐ Tanh
- ☒ **ReLU**
- ☐ Linear

Solution: (C). ReLU addresses the vanishing gradient problem by having a constant gradient of 1 for positive inputs.

3. Attention Mechanism

In the scaled dot-product attention formula $\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}(\mathbf{Q}\mathbf{K}^T / \sqrt{d_k})\mathbf{V}$, what is the purpose of the scaling factor $\sqrt{d_k}$?

- ☐ To normalize the attention weights
- ☒ **To prevent the dot products from becoming too large**
- ☐ To ensure the softmax outputs sum to 1
- ☐ To make the computation more efficient

Solution: (B). The scaling factor prevents the dot products from becoming too large, which would cause the softmax to saturate.

4. GPT Evolution

Which GPT model first demonstrated significant zero-shot capabilities without task-specific fine-tuning?

- ☐ GPT-1

- ☒ GPT-2
- ☐ GPT-3
- ☐ ChatGPT

Solution: (B). GPT-2 was the first to demonstrate significant zero-shot capabilities through scale and prompting.

Written Answer (80 points total)

5. Attention Scaling (20 points)

- Derive why variance of dot product $q \cdot k$ grows with d_k if $q, k \sim \mathcal{N}(0, 1)$.
- Show dividing by $\sqrt{d_k}$ stabilizes variance.
- What happens to attention weights when d_k becomes very large without scaling? Explain the practical implications for training transformers.

Solution:

- a) If $q, k \sim \mathcal{N}(0, I_{d_k})$, then $q \cdot k = \sum_{i=1}^{d_k} q_i k_i$.

Since q_i, k_i are independent standard normal variables, $q_i k_i$ has mean 0 and variance $\text{Var}(q_i) \cdot \text{Var}(k_i) = 1 \cdot 1 = 1$.

Therefore: $\text{Var}(q \cdot k) = \text{Var}\left(\sum_{i=1}^{d_k} q_i k_i\right) = \sum_{i=1}^{d_k} \text{Var}(q_i k_i) = d_k$

- b) Scaling by $\frac{1}{\sqrt{d_k}}$:

$$\text{Var}\left(\frac{q \cdot k}{\sqrt{d_k}}\right) = \frac{1}{d_k} \cdot \text{Var}(q \cdot k) = \frac{1}{d_k} \cdot d_k = 1$$

- c) When d_k becomes very large without scaling:

- **Attention weights become uniform:** As $d_k \rightarrow \infty$, the variance of $q \cdot k$ grows without bound, making all dot products very large
- **Loss of selectivity:** The model can no longer focus on specific relevant tokens, losing the ability to learn meaningful attention patterns

Practical implications for training:

- **Gradient problems:** Uniform attention weights lead to vanishing gradients, making training extremely slow or impossible
- **Training instability:** Large, unnormalized dot products can cause numerical instability and overflow issues

This is why the $\sqrt{d_k}$ scaling is crucial for transformer training - it maintains the model's ability to learn meaningful attention patterns regardless of the embedding dimension.

6. Transformer Architecture and Multi-Head Attention (20 points)

You're implementing a transformer model and need to understand the attention mechanism in detail.

- (7 pts) Explain how multi-head attention works. What are the benefits of having multiple attention heads?
- (7 pts) Given query Q, key K, and value V matrices of shape (seq_len, d_model), derive the computational complexity of self-attention.
- (6 pts) How would you modify the attention mechanism to handle very long sequences (10k+ tokens) efficiently?

Solution:

a) **Multi-head attention:**

- **Process:** Split d_{model} into h heads, each with dimension $d_k = d_{\text{model}}/h$
- **Computation:** Each head computes attention independently: $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$
- **Concatenation:** $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$
- **Benefits:** Captures different types of relationships (syntax, semantics, long-range), parallelizable, more expressive than single head

b) **Computational complexity:**

- **Attention weights:** QK^T has shape $(\text{seq_len}, \text{seq_len})$, complexity $O(\text{seq_len}^2 \cdot d_{\text{model}})$
- **Softmax:** $O(\text{seq_len}^2)$ for each head
- **Value multiplication:** $O(\text{seq_len}^2 \cdot d_{\text{model}})$
- **Total:** $O(\text{seq_len}^2 \cdot d_{\text{model}})$ - quadratic in sequence length

c) **Efficient attention for long sequences:**

- **Sparse attention:** Only attend to nearby tokens or selected positions
- **Sliding window:** Limit attention to a fixed window size
- **Hierarchical attention:** Attend at multiple granularities (sentence, paragraph, document)
- **Linear attention:** Use approximations like Performer or Linformer
- **Memory-efficient:** Gradient checkpointing, recompute attention weights

7. Value Iteration in an MDP (20 points)

Consider a 3-state MDP: (s_1, s_2, s_3) with deterministic transitions and rewards:

- s_1 : $A \rightarrow s_1$, reward 1; $B \rightarrow s_2$, reward 0.
- s_2 : $A \rightarrow s_2$, reward 2; $B \rightarrow s_3$, reward 0.
- s_3 : any action $\rightarrow s_3$, reward 5.

Discount $\gamma = 0.9$.

- (8 pts) Write the Bellman optimality equation.
- (8 pts) Starting with $V_0(s) = 0$, perform one value iteration update. Compute $V_1(s_1)$, $V_1(s_2)$, $V_1(s_3)$.
- (4 pts) Which action is preferred in s_1 after one iteration?

Solution:

- (a) The Bellman optimality equation is:

$$V^*(s) = \max_a \{R(s, a) + \gamma V^*(s')\}$$

where s' is the next state after taking action a in state s .

(b) Starting with $V_0(s) = 0$ for all states, one value iteration update:

For s_1 :

- Action A: $R(s_1, A) + \gamma V_0(s_1) = 1 + 0.9 \cdot 0 = 1$
- Action B: $R(s_1, B) + \gamma V_0(s_2) = 0 + 0.9 \cdot 0 = 0$
- $V_1(s_1) = \max(1, 0) = 1$

For s_2 :

- Action A: $R(s_2, A) + \gamma V_0(s_2) = 2 + 0.9 \cdot 0 = 2$
- Action B: $R(s_2, B) + \gamma V_0(s_3) = 0 + 0.9 \cdot 0 = 0$
- $V_1(s_2) = \max(2, 0) = 2$

For s_3 :

- Any action: $R(s_3, \cdot) + \gamma V_0(s_3) = 5 + 0.9 \cdot 0 = 5$
- $V_1(s_3) = 5$

(c) At s_1 after one iteration, action A is preferred because it gives value 1 vs action B's value 0.

8. GPT Alignment & Safety (20 points)

Your LLM is generating harmful content and you need to implement RLHF to fix this.

- (10 pts) How would you design the human preference data collection process for RLHF?
- (10 pts) Your model is reward hacking (optimizing for the reward model instead of being helpful). How would you detect and fix this?

Solution:

a) Data Collection Process:

- **Prompt design:** Cover helpfulness, safety, factuality across domains and difficulty levels
- **Response generation:** Sample diverse responses using different temperatures/models
- **Annotation:** Pairwise preferences with 5-point rubrics (Helpful, Honest, Harmless)
- **Quality control:** Inter-annotator agreement ($\alpha \geq 0.6$), rater calibration, adversarial checks
- **Sampling:** Active learning on uncertain/disagreed cases, upweight failure modes
- **Bias mitigation:** Track subgroup performance, use counterfactual prompts

b) Reward Hacking Detection & Fix:

- **Detection:** Monitor RM-human gap, metric drift (length/repetition), ensemble disagreement, stress tests
- **Data fixes:** Add hard negatives, mine adversarial examples, include absolute ratings
- **RM fixes:** Calibrate/regularize, use ensemble/pessimistic RM, multi-objective optimization
- **RL fixes:** Increase KL penalty, add exploit penalties, use constrained RL
- **Evaluation:** A/B testing with human judgment, frozen challenge sets, canary monitoring